

## Aims:

This exercise aims to get you to:

1. Practice the use of Elasticsearch
2. Create an Elasticsearch index
3. Perform operations on the created index including queries using both query strings and Elasticsearch's DSL

## Running Elasticsearch Server

ElasticSearch is a NoSQL indexing and search engine built on top of Apache Lucene. It allows you to create indexes over unstructured / semi-structured data / structured data. In order to run Elasticsearch server, first you need to open a terminal (command line tool) in VLab and run the following program:

```
$ 9313
```

This program will setup the environment needed to run Elasticsearch server. You should see an output like the following:

```
$ 9313
Welcome to COMP9313!
newclass starting new subshell for class COMP9313...
```

Next, run the command below to start Elasticsearch server:

```
$ elasticsearch
...
[2019-07-14T17:41:11,470][INFO ][o.e.h.n.Netty4HttpServerTransport] [8Uw8EI3]
publish_address {127.0.0.1:9200}, bound_addresses {127.0.0.1:9200}
[2019-07-14T17:41:11,470][INFO ][o.e.n.Node] [8Uw8EI3] started
[2019-07-14T17:41:11,479][INFO ][o.e.g.GatewayService] [8Uw8EI3] recovered [0] indices
into cluster_state
...
```

If everything goes fine, you should see an output similar to the one above. The server should be able to receive HTTP requests (using Elasticsearch's REST APIs) on <http://localhost:9200>. You can use the command below to check if Elasticsearch is up and running:

```
$ curl "localhost:9200/"

{
  "name" : "8Uw8EI3",
  "cluster_name" : "elasticsearch",
  "cluster_uuid" : "1CEB4Ac1QvqBcdI_VzUXHQ",
  "version" : {
    "number" : "6.1.1",
    "build_hash" : "bd92e7f",
    "build_date" : "2017-12-17T20:23:25.338Z",
    "build_snapshot" : false,
    "lucene_version" : "7.1.0",
    "minimum_wire_compatibility_version" : "5.6.0",
    "minimum_index_compatibility_version" : "5.0.0"
  },
  "tagline" : "You Know, for Search"
}
```

If the server is up and running properly, you should see an output similar to the one shown above. Notice also that two new directories were created in your `HOME` directory, namely, `elasticsearch_data` and `elasticsearch_log`. The first will store all data of the indexes you create in ElasticSearch, while the second one will store logging data.

### Create a new index to store vulnerability-related data

Let's now create a new index to store security vulnerability related data. Open a new terminal and run again the program 9313. Then, run the command below:

```
$ curl -X PUT "http://localhost:9200/vuln_demo?pretty"

{
  "acknowledged" : true,
  "shards_acknowledged" : true,
  "index" : "vuln_demo"
}
```

The command `curl` shown above is a program that allows to send HTTP request to a server. The command above sends a `PUT` HTTP request to ElasticSearch to create an index named `vuln_demo`. If everything goes fine, you should see the output shown above.

You can check the recently created index with the command below:

```
$ curl -X GET "http://localhost:9200/_cat/indices?v"

health status index      uuid                                pri rep docs.count docs.deleted store.size pri.store.size
yellow open    vuln_demo b49GKK3XTMyV7FvSrFyJOQ           5   1         0          0       1.1kb         1.1kb
```

The output shows the created index along with additional information such as the number of replicas (`rep`), document counts (`docs.count`), store size (`store.size`), among other details.

### Create a new mapping

Now, we will create a new mapping within the index just created. This mapping provides a “schema” for the vulnerability data we want to store in our index.

```
$ curl -X PUT "http://localhost:9200/vuln_demo/vulnerability/_mapping?pretty" -H "Content-type: application/json" -d '{"vulnerability": {"properties": {"id": {"type": "text"}, "description": {"type": "text"}, "cvss_score": {"type": "double"}}}}'

{
  "acknowledged" : true
}
```

If the command executes properly, you should see the output `{"acknowledged":true}`.

### Create, modify and delete documents

You can **create** a new document in the document with the command below:

```
$ curl -X PUT "http://localhost:9200/vuln_demo/vulnerability/CVE-1999-0001?pretty" -H "Content-type: application/json" -d '{"id": "CVE-1999-0001", "description": "ip_input.c in BSD-derived TCP/IP implementations allows remote attackers to cause a denial of service (crash or hang) via crafted packets.", "cvss_score": 5.0}'

{
  "index" : "vuln_demo",
  "type" : "vulnerability",
  "_id" : "CVE-1999-0001",
  "_score" : null,
  "_source" : {
    "id" : "CVE-1999-0001",
    "description" : "ip_input.c in BSD-derived TCP/IP implementations allows remote attackers to cause a denial of service (crash or hang) via crafted packets.",
    "cvss_score" : 5.0
  },
  "highlight" : {}
}
```

```
{
  "_type" : "vulnerability",
  "_id" : "CVE-1999-0001",
  "_version" : 1,
  "result" : "created",
  "_shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "_seq_no" : 0,
  "_primary_term" : 1
}
```

You can quickly check (retrieve from index) the recently created document with the command below:

```
$ curl -X GET "http://localhost:9200/vuln_demo/vulnerability/_search?pretty"
{
  "took" : 2,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : 2,
    "max_score" : 1.0,
    "hits" : [
      {
        "_index" : "vuln_demo",
        "_type" : "vulnerability",
        "_id" : "_update",
        "_score" : 1.0,
        "_source" : {
          "cvss_score" : 6.5
        }
      },
      {
        "_index" : "vuln_demo",
        "_type" : "vulnerability",
        "_id" : "CVE-1999-0001",
        "_score" : 1.0,
        "_source" : {
          "cvss_score" : 5.0
        }
      }
    ]
  }
}
```

You can also **update** a document using the following command:

```
$ curl -X POST "http://localhost:9200/vuln_demo/vulnerability/CVE-1999-0001?pretty" -H
"Content-type: application/json" -d '{"cvss_score":6.5}'
{
  "_index" : "vuln_demo",
  "_type" : "vulnerability",
  "_id" : "CVE-1999-0001",
  "_version" : 4,
  "result" : "updated",
  "_shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "_seq_no" : 4,
  "_primary_term" : 1
}
```

```
}
```

To **delete** a document, you can use the command below:

```
$ curl -X DELETE "http://localhost:9200/vuln_demo/vulnerability/CVE-1999-0001?pretty"
{
  "_index" : "vuln_demo",
  "_type" : "vulnerability",
  "_id" : "CVE-1999-0001",
  "_version" : 5,
  "result" : "deleted",
  "shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "_seq_no" : 5,
  "_primary_term" : 1
}
```

## Querying data

Before providing examples on how to query data, let's create a couple of documents:

```
$ curl -X PUT "http://localhost:9200/vuln_demo/vulnerability/CVE-1999-0001?pretty" -H
"Content-type: application/json" -d "{\"id\": \"CVE-1999-0001\", \"description\": \"ip_input.c
in BSD-derived TCP/IP implementations allows remote attackers to cause a denial of service
(crash or hang) via crafted packets.\", \"cvss_score\": 5.0}"

{
  "_index" : "vuln_demo",
  "_type" : "vulnerability",
  "_id" : "CVE-1999-0001",
  "_version" : 3,
  "result" : "created",
  "shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "_seq_no" : 10,
  "_primary_term" : 1
}
```

```
$ curl -X PUT "http://localhost:9200/vuln_demo/vulnerability/CVE-1999-0002?pretty" -H
"Content-type: application/json" -d "{\"id\": \"CVE-1999-0002\", \"description\": \"Buffer
overflow in NFS mountd gives root access to remote attackers, mostly in Linux systems.\",
\"cvss_score\": 10.0}"

{
  "_index" : "vuln_demo",
  "_type" : "vulnerability",
  "_id" : "CVE-1999-0002",
  "_version" : 1,
  "result" : "created",
  "shards" : {
    "total" : 2,
    "successful" : 1,
    "failed" : 0
  },
  "_seq_no" : 0,
  "_primary_term" : 1
}
```

Let's now search for any document containing the string "BSD" anywhere in the document:

```
$ curl -X GET "http://localhost:9200/vuln_demo/vulnerability/_search?pretty&q=BSD"
{
  "took" : 10,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : 1,
    "max_score" : 0.2876821,
    "hits" : [
      {
        "_index" : "vuln_demo",
        "_type" : "vulnerability",
        "_id" : "CVE-1999-0001",
        "_score" : 0.2876821,
        "_source" : {
          "id" : "CVE-1999-0001",
          "description" : "ip_input.c in BSD-derived TCP/IP implementations allows remote
attackers to cause a denial of service (crash or hang) via crafted packets.",
          "cvss_score" : 5.0
        }
      }
    ]
  }
}
```

Next, let's query our index by field and using Boolean operators (notice that use the string "%20" for representing spaces in the URL below):

```
$ curl -X GET
"http://localhost:9200/vuln_demo/vulnerability/_search?pretty&q=description:(TCP%20OR%20BSD)"
{
  "took" : 7,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : 1,
    "max_score" : 0.5753642,
    "hits" : [
      {
        "_index" : "vuln_demo",
        "_type" : "vulnerability",
        "_id" : "CVE-1999-0001",
        "_score" : 0.5753642,
        "_source" : {
          "id" : "CVE-1999-0001",
          "description" : "ip_input.c in BSD-derived TCP/IP implementations allows remote
attackers to cause a denial of service (crash or hang) via crafted packets.",
          "cvss_score" : 5.0
        }
      }
    ]
  }
}
```

Yet, another query including a Boolean expression a little bit more complex than the previous one:

```
$ curl -X GET
"http://localhost:9200/vuln_demo/vulnerability/_search?pretty&q=(description:(Buffer%20AND%
20NFS)%20AND%20cvss_score:>4.0)"

{
  "took" : 7,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : 1,
    "max_score" : 1.5753641,
    "hits" : [
      {
        "_index" : "vuln_demo",
        "_type" : "vulnerability",
        "_id" : "CVE-1999-0002",
        "_score" : 1.5753641,
        "_source" : {
          "id" : "CVE-1999-0002",
          "description" : "Buffer overflow in NFS mountd gives root access to remote
attackers, mostly in Linux systems.",
          "cvss_score" : 10.0
        }
      }
    ]
  }
}
```

Let's now search for vulnerabilities related to "TCP", but let's exclude documents containing "NFS".

```
$ curl -X GET
"http://localhost:9200/vuln_demo/vulnerability/_search?pretty&q=description:%20TCP%20-NFS"

{
  "took" : 6,
  "timed_out" : false,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : 1,
    "max_score" : 0.2876821,
    "hits" : [
      {
        "_index" : "vuln_demo",
        "_type" : "vulnerability",
        "_id" : "CVE-1999-0001",
        "_score" : 0.2876821,
        "_source" : {
          "id" : "CVE-1999-0001",
          "description" : "ip_input.c in BSD-derived TCP/IP implementations allows remote
attackers to cause a denial of service (crash or hang) via crafted packets.",
          "cvss_score" : 5.0
        }
      }
    ]
  }
}
```

### Deleting an index

The command below will delete the index we created for this exercise.

```
$ curl -X DELETE "http://localhost:9200/vuln_demo?pretty"
{
  "acknowledged" : true
}
```

### Stopping Elasticsearch server

You can stop the server with the combination of keys CTRL+C in the terminal where the server is running.